

CONvenient Interface to Inverse Ising (ConIII): A Python package for solving maximum entropy models

¹Edward D Lee, ²Bryan C Daniels

¹*Department of Physics, 142 Sciences Dr, Cornell University, Ithaca NY 14853,*

²*ASU-SFI Center for Biosocial Complex Systems, Arizona State University, Tempe, AZ 85287*

ConIII (pronounced CON-ee) is an open-source Python project providing a simple interface to solving maximum entropy models, with a focus on the Ising model. We describe the maximum entropy problem and give an overview of the algorithms that are implemented as part of ConIII (<https://github.com/bcdaniels/coniii>) including a regularized mean field method, Monte Carlo histogram, pseudolikelihood, and minimum probability flow. We briefly explain how one should approach and validate maxent models from the perspective of model selection. Our goal is to make a variety of maximum entropy techniques accessible to those unfamiliar with the techniques and accelerate workflow for users.

INTRODUCTION

Many biological and social systems are characterized by collective behavior: the correlated pattern of neuron firing [1], protein diversity in the immune system [?], conflict participation in monkeys [?], flocking in birds [2], statistics of letters in words [?], or consensus voting in the US Supreme Court [3]. Statistical physics is a natural approach to probing such systems precisely because they are collective [4]. Recently, the development of numerical, analytic, and computational tools have made it feasible in these large collective systems to solve for the statistical model that reproduces system behavior, corresponding to solving an “inverse problem.” This approach contrasts with the typical problem in statistical physics where one postulates the microscopic model (the Hamiltonian) and works out the physical behavior of the system. In the inverse problem, we find the parameters that correspond to observed behavior of a known system. In many cases, this is a very difficult problem to solve and does not have an analytical solution, and we must rely on analytic approximation and numerical techniques to estimate the parameters.

The Ising model has been of particular interest because of its simplicity and generality. A variety of algorithms have been proposed to solve the inverse Ising problem, but different approaches are disparately available on separate code bases in different coding languages, which makes comparison difficult and pedagogy more complicated. ConIII (Convenient Interface to Inverse Ising) is a Python project intended to provide a centralized resource for the inverse Ising problem and provide a base for the addition of more maximum entropy problems in the future. With ConIII, it is possible to solve the inverse Ising problem with a variety of algorithms in just a few lines of code.

WHAT IS MAXIMUM ENTROPY?

Shannon introduced the concept of information entropy in his seminal paper about communication over a noisy channel [5]. Information entropy is the unique measure of uncertainty that follows from insisting on some elementary principles of consistency. According to Shannon, the entropy over the probability distribution $p(s)$ of possible discrete configurations \mathcal{S} of a system is

$$S[p] = - \sum_{s \in \mathcal{S}} p(s) \log p(s). \quad (1)$$

These configurations could be on-off patterns of firing in neurons, the arrangement of letters in a word, or the orientation of spins in a material.

When there is no structure in the distribution, meaning that the probability is uniform, entropy is at a maximum. In the context of communication theory as Shannon first discussed, this means that there is no structure to exploit to make a prediction about the next part of an incoming message; thus, maximum entropy means that each new part of the message is maximally “surprising.” At the other extreme, when the message consists of the same bit over and over again, we can always guess at the following part of the message and the signal has zero entropy. In the context of modeling, we use entropy not to refer to the difficulty of the message, but to our state of knowledge about it. Entropy precisely measures our uncertainty about which configuration we will find the system to be in.

Maximum entropy, or maxent, is the formal framework for building models that are consistent with statistics from the data but otherwise as structureless as possible [6, 7]. We begin by choosing that there is a set of useful or important features from the data $f_k(s)$ that should be true for the model that we are trying to build. These could be the averaging firing probabilities of a set of neurons or the average pairwise cooccurrence for primates in a sequence of conflicts. The average of this feature across

set of states in the data set \mathcal{D} is

$$\langle f_k \rangle_{\text{data}} = \frac{1}{R} \sum_{s \in \mathcal{D}} f_k(s) \quad (2)$$

According to the model in which each observation s occurs with some probability $p(s)$, the same average is calculated over all possible states

$$\langle f_k \rangle = \sum_{s \in \mathcal{S}} p(s) f_k(s) \quad (3)$$

We assert that the model should fit the K features while maximizing entropy. The standard procedure is to solve this by the method of Lagrangian multipliers. We construct the Lagrangian functional \mathcal{L} by introducing the multipliers λ_k .

$$\mathcal{L}[p] = - \sum_{s \in \mathcal{S}} p(s) \log p(s) - \sum_{k=1}^K \lambda_k (\langle f_k \rangle - \langle f_k \rangle_{\text{data}}) \quad (4)$$

Then, we solve for the fixed point by taking the derivative with respect to λ_k . The resulting model is a Boltzmann distribution over states:

$$p_{\text{ME}}(s) = e^{-E(s)} / Z, \quad (5)$$

with relative negative log-likelihood (also known as the energy or Hamiltonian)

$$E(s) = - \sum_{k=1}^K \lambda_k f_k(s), \quad (6)$$

and normalization factor (also known as the partition function)

$$Z = \sum_{s \in \mathcal{S}} e^{-E(s)}. \quad (7)$$

Entropy is a convex function of p and the constraints are linear with respect to p , so the problem is convex and the maxent distribution unique. Readers familiar with statistical physics will recognize this as an alternative derivation of the microcanonical ensemble, demonstrating that statistical mechanics can be viewed as an inference procedure using the maximum entropy principle [7].

Finding the parameters λ_k that match the constraints $\langle f_k \rangle_{\text{data}}$ is equivalent to minimizing the Kullback-Leibler divergence between the model and the data [8]

$$D_{\text{KL}}(p_{\text{data}} || p_{\text{ME}}) = \sum_s p_{\text{data}} \log \left(\frac{p_{\text{data}}(s)}{p_{\text{ME}}(s)} \right) \quad (8)$$

$$\frac{\partial D_{\text{KL}}}{\partial \lambda_k} = \sum_s p_{\text{data}}(s) \frac{\partial (-E(s) - \log Z)}{\partial \lambda_k} = 0$$

$$\Rightarrow \langle f_k \rangle_{\text{data}} = \langle f_k \rangle_{\text{ME}}. \quad (9)$$




FIG. 1. Example of a fully connected Ising model with random couplings. Each spin s_i (circle) can take one of two states (black or white, corresponding to -1 and $+1$) and is connected to every other spin in the system with a positive or negative coupling. This schematic represents one possible state of the spins. The inverse Ising problem starts with statistics about the states observed in the system and infers the coupling strengths J_{ij} and individual biases h_i that produce such statistics.

In other words, the parameters of the maximum entropy model are the ones that minimize the information theoretic “distance” to the distribution of the data. Note that these parameters are given by the data: once the constraints have been chosen, there is a single maximum entropy solution, with no free parameters.

The Ising model

The Ising model is a statistical physics model of magnetism [?]. It consists of a set of spins s_i with 2 possible orientations (up and down), each coupled to an external magnetic field h_i and coupled to each other with couplings J_{ij} . The strength of the magnetic field determines the tendency of each of the spins to orient in a particular direction and the couplings J_{ij} determine whether the spins tend to point together ($J_{ij} > 0$) or against each other ($J_{ij} < 0$). Typically, neighbors are defined as spins that interact with one another given by some underlying network structure: Figure 1 shows a fully-connected example.

The energy of each configuration determines its probability via (5):

$$E = - \sum_{i < j}^N J_{ij} s_i s_j - \sum_{i=1}^N h_i s_i. \quad (10)$$

We can derive the Ising model from the perspective of

maximum entropy. Fixing the means and pairwise correlations to those observed in the data

$$\langle s_i \rangle = \langle s_i \rangle_{\text{data}} \quad (11)$$

$$\langle s_i s_j \rangle = \langle s_i s_j \rangle_{\text{data}} \quad (12)$$

we go through the procedure of constructing the Lagrangian from Eq 4

$$\mathcal{L}[p] = - \sum_s p(s) \log p(s) + \sum_{i < j}^N J_{ij} \langle s_i s_j \rangle + \sum_i^N h_i \langle s_i \rangle \quad (13)$$

$$\frac{\partial \mathcal{L}[p]}{\partial p(s)} = -\log p(s) - 1 + \sum_{i < j}^N J_{ij} s_i s_j + \sum_i^N h_i s_i \quad (14)$$

$$\log p(s) = -1 + \sum_{i < j}^N J_{ij} s_i s_j + \sum_i^N h_i s_i \quad (15)$$

$$p(s) = e^{-E(s)} / Z \quad (16)$$

where to enforce normalization of the probability distribution $\sum_s p(s) = 1$,

$$Z = \sum_s e^{-E(s)}. \quad (17)$$

Thus, the resulting model is exactly the Ising model mentioned earlier.

Despite the simplicity of the Ising model, the structure imposed by the discrete nature of the spins means that finding the parameters is challenging analytically and computationally. In the last few years, numerous techniques have been suggested for solving the inverse Ising problem exactly or approximately [9]. We have implemented some of them in ConIII and designed a package structure to make it easily extensible to include more methods. Here, we briefly describe the algorithms that are part of the first official version of the package. The goal is to give the user a sense (or reminder) of how they work without getting bogged down in heavy detail. For more detail, we suggest perusing the papers referenced in each section or the review [9]. For a complete beginner, it may be useful to first get familiar with a slower introduction like in the Appendix of Ref [3], Ref [?], or Ref [6].

INVERSE ISING METHODS IMPLEMENTED IN CONIII

Enumeration

The naïve approach that only works for small systems is to write out the equations from Eq 9 and solve them numerically. After writing out all K equations,

$$\langle f_k \rangle = - \frac{\partial \ln Z}{\partial \lambda_k} = \langle f_k \rangle_{\text{data}}, \quad (18)$$

we can use any standard root-finding algorithm to find the parameters λ_k . This approach, however, involves enumerating all terms in the partition function Z , whose number grows exponentially with system size.

For the Ising model, writing down the equations has a number of steps $\mathcal{O}(K^2 2^N)$, where K is the number of constraints and N the number of spins. Each evaluation of the objective in the root-finding algorithm will be of the same order. For relatively small systems, around $N \leq 15$, this approach is feasible on a typical desktop computer and is a good way to test the results of a more complicated algorithm.

This approach is part of the `Exact` class that contains code for writing Eqs 9 into a file and solving them with the `scipy.optimize` library.

Monte Carlo method

Perhaps the most straightforward and most expensive computational approach is Monte Carlo Markov Chain (MCMC) sampling. A series of states sampled from a proposed $p(s)$ is produced by MCMC to approximate $\langle f_k \rangle$ and determine how close we are to matching $\langle f_k \rangle_{\text{data}}$. The parameters are then adjusted using a learning rule, and both sampling and learning are repeated until a stopping criterion is met. This can be combined with a variety of approximate gradient descent methods to reduce the number of sampling steps by predicting how the distribution will change if we modify the parameters slightly. The particular technique implemented in ConIII is the Monte Carlo Histogram (MCH) method [10].

Since the sampling step is expensive, the idea behind MCH is to reuse a sample for more than one gradient descent step [10]. Given that we have a sample with probability distribution $p(s)$ generated with parameters λ_k , we would like to estimate the proposed distribution $p'(s)$ from adjusting our parameters $\lambda'_k = \lambda_k + \Delta \lambda_k$. We can leverage our current sample to make this extrapolation.

$$p' = \frac{p'}{p} p \quad (19)$$

$$p'(s) = \frac{Z}{Z'} e^{\sum_k \Delta \lambda_k f_k(s)} p(s) \quad (20)$$

To estimate the average,

$$\sum_s p'(s) f_k(s) = \frac{Z}{Z'} \sum_s p(s) e^{\sum_k \Delta \lambda_k f_k(s)} f_k(s) \quad (21)$$

To be explicit about the fact that we only have a sampled approximation to p , we replace p with the sample distribution.

$$\langle f_k \rangle' = \frac{Z}{Z'} \left\langle e^{\sum_k \Delta \lambda_k f_k(s)} f_k(s) \right\rangle_{\text{sample}} \quad (22)$$

Likewise, the ratio of the partition function can be estimated

$$\frac{Z}{Z'} \approx 1 \left/ \left\langle e^{\sum_k \Delta \lambda_k f_k(s)} \right\rangle_{\text{sample}} \right. \quad (23)$$

At each step, we update the Lagrangian multipliers $\{\lambda_k\}$ while being careful to stay within the bounds of a reasonable extrapolation. One suggestion is to update the parameters with some inertia [?]]

$$\Delta \lambda_k(t+1) = \Delta \lambda_k(t) + \epsilon \Delta \lambda_k(t-1) \quad (24)$$

$$\Delta \lambda_k(t) = \eta (\langle f_k \rangle' - \langle f_k \rangle) \quad (25)$$

This has a fixed point at the correct parameters.

In practice, MCH can be difficult to tune properly and one must check in on the progress of the algorithm often. One issue is choosing how to set the learning rule parameters η and ϵ . One suggestion for η is to shrink it as the inverse of the number of iterations [?]]. Another issue is that parameters cannot be changed by too much when using the MCH approximation step or the extrapolation to λ'_k will be inaccurate and the algorithm will fail to converge. In ConIII, this can be controlled by setting a bound on the maximum possible change in each parameter $\Delta \lambda_{\max}$ and restricting the norm of the vector of change in parameters $\sum_k \sqrt{\Delta \lambda_k^2}$. Another issue is setting the parameters of the MCMC sampling routine. Both the burn time (the number of iterations before starting to sample) and sampling iterations (number of iterations between samples) must be large enough that we are sampling from the equilibrium distribution. Typically, these are found by measuring how long the energy or individual parameter values remain correlated as MCMC progresses. The parameters may need to be updated during the course of MCH because the sampling parameters may need to change with the estimated parameters of the model. For some regimes of parameter space, samples are correlated over long times and alternative sampling methods like Wolff or Swendsen-Wang should be used. We do not discuss these sampling details here, but see Refs [11?]] for examples.

The main computation cost for MCH lies in the sampling step. For each iteration of MCH, the runtime is proportional to the number of samples n_{sample} , number of MCMC iterations T_{MC} , the number of constraints for the Ising model N^2 : $\mathcal{O}(T_{\text{MC}} n_{\text{sample}} N^2)$, whereas the MCH estimate is relatively quick $\mathcal{O}(n_{\text{MCH}} n_{\text{sample}} N^2)$ because the number of MCH approximation steps needed to converge is much smaller than the number of MCMC sampling iterations $n_{\text{MCH}} \ll T_{\text{MC}}$.

MCH is implemented in the `MCH` class.

Pseudolikelihood

The pseudolikelihood approach is an analytic approximation to the likelihood that drastically reduces the com-

putational complexity of the problem and is exact as $N \rightarrow \infty$ [12]. We calculate the conditional probability of each spin s_i given the rest of the system $\{s_{j \neq i}\}$

$$p(s_i | \{s_{j \neq i}\}) = \left(1 + e^{-2s_i (h_i + \sum_{j \neq i} J_{ij} s_j)} \right)^{-1} \quad (26)$$

Taking the logarithm, we define the approximate log-likelihood by summing over data points indexed by r :

$$f(h_i, \{J_{ij}\}) = \sum_{r=1}^R \ln p(s_i^{(r)} | s_{\setminus i}^{(r)}) \quad (27)$$

In the limit where the ensemble is well sampled, the average over the data can be replaced by an average over the ensemble:

$$f(h_i, \{J_{ij}\}) = \sum_s \ln p(s_i | \{s_{j \neq i}\}) p(s; h_i, \{J_{ij}\}) \quad (28)$$

At maximum likelihood,

$$\frac{\partial f}{\partial J_{ij}} = 0 \quad (29)$$

At iteration through i , the maximum likelihood parameters may not satisfy the symmetry $J_{ij} = J_{ji}$. We impose the symmetry by insisting that at each step the new couplings are

$$J'_{ij} = (J_{ij} + J_{ji})/2 \quad (30)$$

Pseudolikelihood is extremely fast. Each calculation of the gradient of the pseudolikelihood across all spins for the Ising model has runtime $\mathcal{O}(RN^3)$ and is often surprisingly accurate.

Pseudolikelihood for the Ising model is implemented in `Pseudo`.

Minimum Probability Flow

Minimum probability flow involves analytically approximating how the probability distribution *changes* as we modify the *configurations* [13?]]. In the methods so far mentioned, the approach has been to maximize the objective (the likelihood function) by immediately taking the derivative with respect to the parameters. With MPF, we first posit a set of dynamics that will lead the data distribution to equilibrate to that of the model. When these distributions are equivalent, then there is no “probability flow” between them. This technique is closely related to score matching, where we instead have a continuous state space and can directly take the derivative with respect to the states without specifying dynamics [14].

First note that Monte Carlo dynamics (satisfying ergodicity and detailed balance) would lead to equilibration of the two distributions. One such transition matrix suggested in Ref [13] is

$$\dot{p}_s = \sum_{s' \neq s} \Gamma_{ss'} p_{s'} - \sum_{s' \neq s} \Gamma_{s's} p_s \quad (31)$$

$$\Gamma_{ss'} = g_{ss'} \exp \left[\frac{1}{2} (E_{s'} - E_s) \right] \quad (32)$$

with transition probabilities $\Gamma_{ss'}$ from state s' to state s . The connectivity matrix $g_{ss'}$ specifies whether there is edge between states s and s' such that probability can flow between them. When data and model distributions are different, probability will flow between them and indicates that the parameters must be changed. By choosing a sparse $g_{ss'}$ while not breaking ergodicity, we can drastically reduce the computational cost of computing this matrix.

We now minimize a derivative of the Kullback-Leibler divergence, measuring how the difference between the model and the states in the data \mathcal{D} changes when the dynamics are run for an infinitesimal amount of time ϵ .

$$\partial_t D_{\text{KL}}(p^{(0)} || p^{(t)}(\{\lambda_k\})) = \sum_{s \notin \mathcal{D}} \dot{p}_s(\lambda_k) \quad (33)$$

$$L(\{\lambda_k\}) = \sum_{s \notin \mathcal{D}} \dot{p}_s(\lambda_k) \quad (34)$$

The idea is that this derivative is also minimized with optimal parameters: the MPF algorithm looks for a minimum of the objective function L .

For the Ising model, each evaluation of the objective function where $\Gamma_{ss'}$ connects each data state with G neighbors has runtime $\mathcal{O}(RGN^2)$. In a large fully connected system, $G \sim 2^N$ would be prohibitively large so a sparse choice is necessary.

MPF is implemented in the `MPF` class.

Regularized mean-field method

One attractively simple and efficient approach uses a regularized version of mean-field theory. In the inverse Ising problem, mean-field theory is equivalent to treating each binary individual as instead having a continuously varying state (corresponding to its mean value). The inverse problem then turns into simply inverting the correlation matrix C [15]:

$$J_{ij}^{\text{mean-field}} = - \frac{(C^{-1})_{ij}}{\sqrt{p_i(1-p_i)p_j(1-p_j)}}, \quad (35)$$

where

$$C_{ij} = \frac{p_{ij} - p_i p_j}{\sqrt{p_i(1-p_i)p_j(1-p_j)}}, \quad (36)$$

and where p_i corresponds to the frequency of individual i being in the active (+1) state and p_{ij} is the frequency of the pair i and j being simultaneously in the active state.

A simple regularization scheme in this case is to discourage large values in the interaction matrix J . This corresponds to putting more weight on solutions that are closer to the case with no interactions (independent individuals). A particularly convenient form adds the following term, quadratic in J , to the negative log-likelihood:

$$\gamma \sum_i \sum_{j>i} J_{ij}^2 p_i(1-p_i)p_j(1-p_j). \quad (37)$$

In this case, the regularized version of the mean-field solution in (35) can be solved analytically, with the slowest computational step coming from the inversion of the correlation matrix. For details, see Refs. [16?].

The idea is then to vary the regularization strength γ to move between the non-interacting case ($\gamma \rightarrow \infty$) and the naively calculated mean-field solution (35) ($\gamma \rightarrow 0$). While there is no guarantee that varying this one parameter will produce solutions that are good enough to “fit within error bars,” this approach has been successful in at least one case of fitting social interactions [?].

This is implemented in `RegularizedMeanField`.

Cluster expansion

Adaptive cluster expansion [15–17] iteratively calculates terms in the cluster expansion of the entropy S :

$$S = \sum_{\Gamma} \Delta S_{\Gamma}, \quad (38)$$

where the sum is over clusters Γ and in the exact case includes all $2^N - 1$ possible nonempty subsets of individuals in the system.¹ The inverse Ising problem is solved independently on each of the clusters, which can be done exactly when the clusters are small. These results are used to construct a full interaction matrix J_{ij} . The expansion starts with small clusters and expands to use larger clusters, neglecting any clusters whose contribution ΔS_{Γ} to the entropy falls below a threshold. To find the best solution that does not overfit, the threshold is initially set at a large value and then lowered, gradually including more clusters in the expansion, until samples from the resulting J_{ij} fit the desired statistics of the data sufficiently well.

The adaptive cluster expansion method is implemented in the `ClusterExpansion` class.

¹ In the simplest version of the expansion, one expands around $S = 0$. In some cases it can be more advantageous to write the expansion around $S - S_0$, where S_0 is a reference entropy corresponding to an easily calculated case such as the independent individual solution or one of the mean-field solutions described in the previous section [16].

SAMPLERS

For estimating statistics of a given probability distribution $p(s)$, we have implemented in ConIII two versions of the well-known Metropolis algorithm. One is specific to the Ising model (MCIsing) and the other (MC) can sample a system given an arbitrary function that calculates the energy of each state.

In some parameter regimes where spins are tightly correlated, the Metropolis algorithm is very inefficient. Cluster sampling like Wolff or Swendsen-Wang are much more efficient and will be included in future releases of ConIII.

THE FINITE SAMPLE PROBLEM

A fundamental problem in model inference is that uncertainty coming from the finiteness of data translates into uncertainties in parameters. In the exposition of the maxent formulation derived from Eq 4, there is no ambiguity in the model parameters because they are fully specified by the constraints calculated from data. In reality, these constraints are typically estimated from a finite sample, and they are noisy. The straightforward answer to this problem is to take more data—in a pairwise maximum entropy problem, we might insist that we have enough samples to well-constrain the correlation between every pair of individuals. But it is not always possible to take enough data. For instance, in a social system in which we are trying to measure stable social structure that lasts on the order of months, there are only a finite number of social interactions that occur over those months, which may not be enough to tightly constrain parameters. When we fit exactly to constraints measured from a small amount of data, we run the danger of overfitting and poor generalization.

When we find the parameters that minimize the Kullback-Leibler divergence between the model and the data distributions, we are maximizing the likelihood of the data. Sample size influences curvature of the likelihood function around the maximum-likelihood peak. Specifically, small sample size leads to a smaller curvature, implying a broad peak and larger uncertainty in parameters. We can estimate this uncertainty from the data: Assuming that the data is independently and identically distributed, the errors are given by the standard error of the mean of a binomial distribution $p_{ij\dots k} = p(s_i = s_j = \dots = s_k = 1)$ with R data points.

$$\delta_{ij\dots k} = \sqrt{(1 - p_{ij\dots k})p_{ij\dots k}/R}. \quad (39)$$

In algorithms that iteratively bring the model closer to the data, we can stop once we have gotten close enough, where close enough is given by the expected fluctuations

in the data distribution.²

An alternative approach is to regularize the problem as in the regularized mean field or cluster expansion algorithms. Here, we restrict the search space in some principled way so that more complicated solutions are disallowed. We then check that the regularized solutions fit the data within expected statistical fluctuations. If not, a more lax regularization can be used to allow more complicated solutions that are able to fit the remaining signal in the data. In a Bayesian formulation, this approach is equivalent to including a prior distribution over parameters that more heavily weights simpler models.

MODEL VALIDATION

Besides just fitting the model while accounting for the noisiness of the data, how do we know if we need a more complicated model? To answer this question, we need to quantify how well the model fits and the corresponding tradeoff between the complexity and fit of the model.

One way to formalize this is to quantify the tradeoff between a good description of the data with the cost of describing the model. This tradeoff manifests as a competition between the likelihood and prior on the model's parameters in the Bayesian formulation. This picture is formalized by quantities like the Akaike information criterion (AIC) or the Bayesian information criterion (BIC) [?].

Although likelihood tells us how well the models are doing relative to one another, it does not tell us in detail how the model is fitting the data. Basic checks would be to compare against other correlations in the data. This could include higher-order correlations than those used to fit the model or other features that are relevant to the question at hand. For instance, an important coarse-grained feature in macaque conflict is the distribution of fight sizes, which is well reproduced by fitting only pairwise correlations [?]. The most strict check would be to compare the entire model probability distribution with that of the data [3], but this is only feasible when the data set is reasonably large.

A summary of how well the model captures the distribution across the entire probability distribution is the multi-information [1?]. If we take a maxent model and add further constraints, the models can be ordered in terms of entropy $S_1 > S_2 > \dots S_m > \dots > S_{\text{data}}$, where the minimum entropy the most constrained model could have is equal to the data. To measure how much correlation in the data our model has captured, we can calculate the multi-information $I_m = S_1 - S_m$, the amount of correlation captured by the model relative to the independent model. The fraction of multi-information captured

² This is like early stopping [?].

is $F = I_m/I_{\text{data}}$. This is a measure of how much of the correlation in the data is captured by the model.³

These tools to quantify model fit then inform the important choice of which constraints to impose. Typically, the approach is to constrain the lowest order interactions that are sufficient to produce collective behavior. The intuition from physics is from the observation that many physical systems are very well described by pairwise interactions [?].⁴ In fact, just the observation that we can model a system well with only pairwise interactions may be surprising. From the model fitting perspective, however, we might choose to constrain other parts of the probability distribution. Ref [?] explores choosing correlations to constrain depending on whether or not they are significantly large. Ref [?] discusses how a pairwise model becomes an increasingly effective description of a system with higher order interactions as the system gets larger.

CONCLUSION

Science is a process of model selection. In the ideal picture, we start with the simplest models and the fewest constraints possible, and then we increase the complexity of the model til it is sufficiently so to make good predictions. In principle, we could add as many constraints as would allow us to fit the data well, but complex models are more computationally expensive and often do not generalize well. Maxent is a principled framework for this picture of model building. The number of parameters and the order of the constraints we impose can be adjusted to test our hypotheses about what matters most for recreating a system's behavior. In this sense, the maxent approach is a useful framework for statistical inference problems far beyond statistical physics. We have built ConIII as an easily accessible software package in the hopes that it will inspire those unfamiliar with maxent approaches to experiment and apply this technique to their own research questions.

[1] E Schneidman, M J Berry, RS II, and W Bialek. Weak pairwise correlations imply strongly correlated network

states in a neural population. *Nature*, 2006.

- [2] William Bialek, Andrea Cavagna, Irene Giardina, Thierry Mora, Edmondo Silvestri, Massimiliano Viale, and Aleksandra M Walczak. Statistical mechanics for natural flocks of birds. *PNAS*, 109(13):4786–4791, 2012.
- [3] Edward D Lee, Chase P Broedersz, and William Bialek. Statistical Mechanics of the US Supreme Court. *J Stat Phys*, pages 1–27, April 2015.
- [4] Claudio Castellano, Santo Fortunato, and Vittorio Loreto. Statistical physics of social dynamics. *Rev. Mod. Phys.*, 81(2):591–646, May 2009.
- [5] Claude Elwood Shannon. A Mathematical Theory of Communication. *The Bell System Technical Journal*, 27:379–423, July 1948.
- [6] E T Jaynes. *Probability Theory: The Logic of Science*. Cambridge University Press, 2003.
- [7] E T Jaynes. Information Theory and Statistical Mechanics. *Phys. Rev.*, 106(4):620, May 1957.
- [8] Thomas M Cover and Joy A Thomas. *Elements of Information Theory*. 2006.
- [9] H Chau Nguyen, Riccardo Zecchina, and Johannes Berg. Inverse statistical problems: from the inverse Ising problem to data science. *arXiv*, February 2017.
- [10] Tamara Broderick, Miroslav Dudik, Gašper Tkačik, Robert E Schapire, and William Bialek. Faster solutions of the inverse pairwise Ising problem. *arXiv*, pages 1–8, December 2007.
- [11] David J C MacKay. *Information Theory, Inference and Learning Algorithms*, September 2005.
- [12] Erik Aurell and Magnus Ekeberg. Inverse Ising Inference Using All the Data. *Physical Review Letters*, 108(9):090201, March 2012.
- [13] Jascha Sohl-Dickstein, Peter B Battaglino, and Michael R DeWeese. New Method for Parameter Estimation in Probabilistic Models: Minimum Probability Flow. *Physical Review Letters*, 107(22):220601, November 2011.
- [14] Aapo Hyvärinen. Some extensions of score matching. *Computational Statistics & Data Analysis*, 51(5):2499–2512, February 2007.
- [15] S Cocco and R Monasson. Adaptive cluster expansion for the inverse Ising problem: convergence, algorithm and tests. *Journal of Statistical Physics*, pages 1–57, 2012.
- [16] John Barton and Simona Cocco. Ising models for neural activity inferred via selective cluster expansion: structural and coding properties. *Journal of Statistical Mechanics: Theory and Experiment*, 2013(03):P03002, mar 2013.
- [17] S Cocco and R Monasson. Adaptive Cluster Expansion for Inferring Boltzmann Machines with Noisy Data. *Physical Review Letters*, 106(9):090601, March 2011.

³ See Refs [?] and [3] for details and further references on how to estimate information quantities.

⁴ In statistical physics, highly correlated behavior across space and

time can emerge from pairwise interactions and renormalization group flow ensures they are the only relevant parameters in the thermodynamic limit [].